











Review Test Submission: HDFS Quiz

User	Venkata Kowsik Temididapathi
Course	CS 6350.002 - Big Data Management and Analytics - S22
Test	HDFS Quiz
Started	2/2/22 10:10 PM
Submitted	2/2/22 10:27 PM
Due Date	2/6/22 11:59 PM
Status	Completed
Attempt Score	100 out of 100 points
Time Elapsed	16 minutes
Results Displayed	All Answers, Submitted Answers, Correct Answers

Question 1

10 out of 10 points

Which of the following are features of HDFS?

- Selected Answers:  It uses inexpensive, commodity hardware
-  It is a multi-node system, rather than a multi-core system.
-  It is fault-tolerant and uses redundancy for storage.
-  It is best suited for batch jobs and can provide high throughput of data processing
- Answers:  It is able to provide real-time, instantaneous processing.
-  It uses inexpensive, commodity hardware
-  It is a multi-node system, rather than a multi-core system.
-  It is a multi-core system, rather than a multi-node system.
-  It is fault-tolerant and uses redundancy for storage.
-  It is best suited for batch jobs and can provide high throughput of data processing


Question 2

10 out of 10 points

HDFS is based on which type of system architecture?

Selected Answer:  Master-Slave

Answers: Peer-to-Peer

 Master-Slave


It has no architecture

Any node can become master for any request.

Question 3

10 out of 10 points

Which one of the following stores data?

Selected Answer:  Datanode

Answers: Namenode

 Datanode

Masternode


Slavenode

Question 4

10 out of 10 points

Which of the following is true if the replication factor is 3

Selected Answers:  File is divided into blocks and each block is written 3 times in HDFS

 The 3 replicas are placed on 3 different nodes

Answers:  File is divided into blocks and each block is written 3 times in HDFS

3 copies of the file are made and they are mirrored 3 times in HDFS

 The 3 replicas are placed on 3 different nodes

The 3 replicas are placed on the same node









Question 5

10 out of 10 points

Match the hadoop shell commands on the left to their meaning on the right

(You can consult this page for help:

<https://hadoop.apache.org/docs/r2.7.1/hadoop-project-dist/hadoop-common/FileSystemShell.html>)

Question	Correct Match	Selected Match
hdfs dfs -copyFromLocal (or hadoop fs -copyFromLocal)	 B. copies a local file to HDFS	 B. copies a local file to HDFS
hdfs dfs -find (or hadoop fs -find)	 A. Used to search for files that match a specified pattern	 A. Used to search for files that match a specified pattern
hdfs dfs -get (or hadoop fs -get)	 C. Copies a HDFS file to local directory	 C. Copies a HDFS file to local directory
hdfs dfs -getfattr (or hadoop fs -getfattr)	 D. Gets extended attributes of the file	 D. Gets extended attributes of the file


All Answer Choices


- A. Used to search for files that match a specified pattern
- B. copies a local file to HDFS
- C. Copies a HDFS file to local directory
- D. Gets extended attributes of the file

Question 6

10 out of 10 points

The client reading the data from HDFS filesystem in Hadoop does which of the following?

Selected Answer:  Gets only the block locations form the namenode


- Answers:
- Gets the data from the namenode
 - Gets the block location from the datanode
 - Gets both the data and block location from the namenode
 -  Gets only the block locations form the namenode

Question 7

10 out of 10 points

In a HDFS system with block size 64MB we store a file which is less than 64MB.Which of the following is true?

Selected Answer:  The file will consume less than 64MB









- Answers:
- The file will consume 64MB
 - The file will consume more than 64MB
 -  The file will consume less than 64MB

It depends on the type of data.

Question 8

10 out of 10 points

Which of the following are true about the namenode data storage and recovery operations?

- Selected Answers:  NN stores all meta-data about block-node mapping in memory
 fsimage is a point-in-time snapshot of the metadata
 Incremental edits after the point-in-time snapshot are written to the edits file
 Secondary namenode periodically merges the fsimage and edits file into a new fsimage file.
- Answers:  NN stores all meta-data about block-node mapping in memory
 fsimage is a point-in-time snapshot of the metadata
 Incremental edits after the point-in-time snapshot are written to the edits file
 Secondary namenode periodically merges the fsimage and edits file into a new fsimage file.



Question 9

10 out of 10 points

Why is the block size in HDFS much larger as compared to other operating systems (e.g. Unix)?

Hint:



<http://fibrevillage.com/storage/645-hadoop-hdfs-blocks-why-is-a-block-in-hdfs-so-large>

- Selected Answer:  To minimize disk seek time as compared to data transfer time
- Answers: To improve network speed
To minimize downtime
To improve disk transfer rate
 To minimize disk seek time as compared to data transfer time

Question 10

10 out of 10 points

As compared to a Relational Database Management System (RDBMS), Hadoop has which of the following properties:

- Selected Answer:  Can work better with structured or unstructured data
- Answers: Has the concept of ACID transactions
Is designed for simultaneous reads and writes
 Can work better with structured or unstructured data

Better data integrity

Saturday, March 12, 2022 9:18:48 AM CST

← OK