

Information Retrieval Processing with MapReduce

Based on Jimmy Lin's Tutorial at the 32nd Annual International ACM SIGIR Conference on Research and Development in Information Retrieval (SIGIR 2009)



This work is licensed under a Creative Commons Attribution-Noncommercial-Share Alike 3.0 United States See <http://creativecommons.org/licenses/by-nc-sa/3.0/us/> for details. PageRank slides adapted from slides by Christophe Bisciglia, Aaron Kimball, & Sierra Michels-Slettvet, Google Distributed Computing Seminar, 2007 (licensed under Creative Commons Attribution 3.0 License)

How much data?

- Google processes 20 PB a day (2008)
- Wayback Machine has 3 PB + 100 TB/month (3/2009)
- Facebook has 2.5 PB of user data + 15 TB/day (4/2009)
- eBay has 6.5 PB of user data + 50 TB/day (5/2009)
- CERN's LHC will generate 15 PB a year (??)



640K ought to be
enough for anybody.

MapReduce

e.g., Amazon Web Services

cheap commodity clusters (or utility computing)
+ simple distributed programming models
+ availability of large datasets
= data-intensive IR research for the masses!

ClueWeb09

ClueWeb09

- NSF-funded project, led by Jamie Callan (CMU/LTI)
- It's big!
 - 1 billion web pages crawled in Jan./Feb. 2009
 - 10 languages, 500 million pages in English
 - 5 TB compressed, 25 uncompressed
- It's available!
 - Available to the research community
 - Test collection coming (TREC 2009)

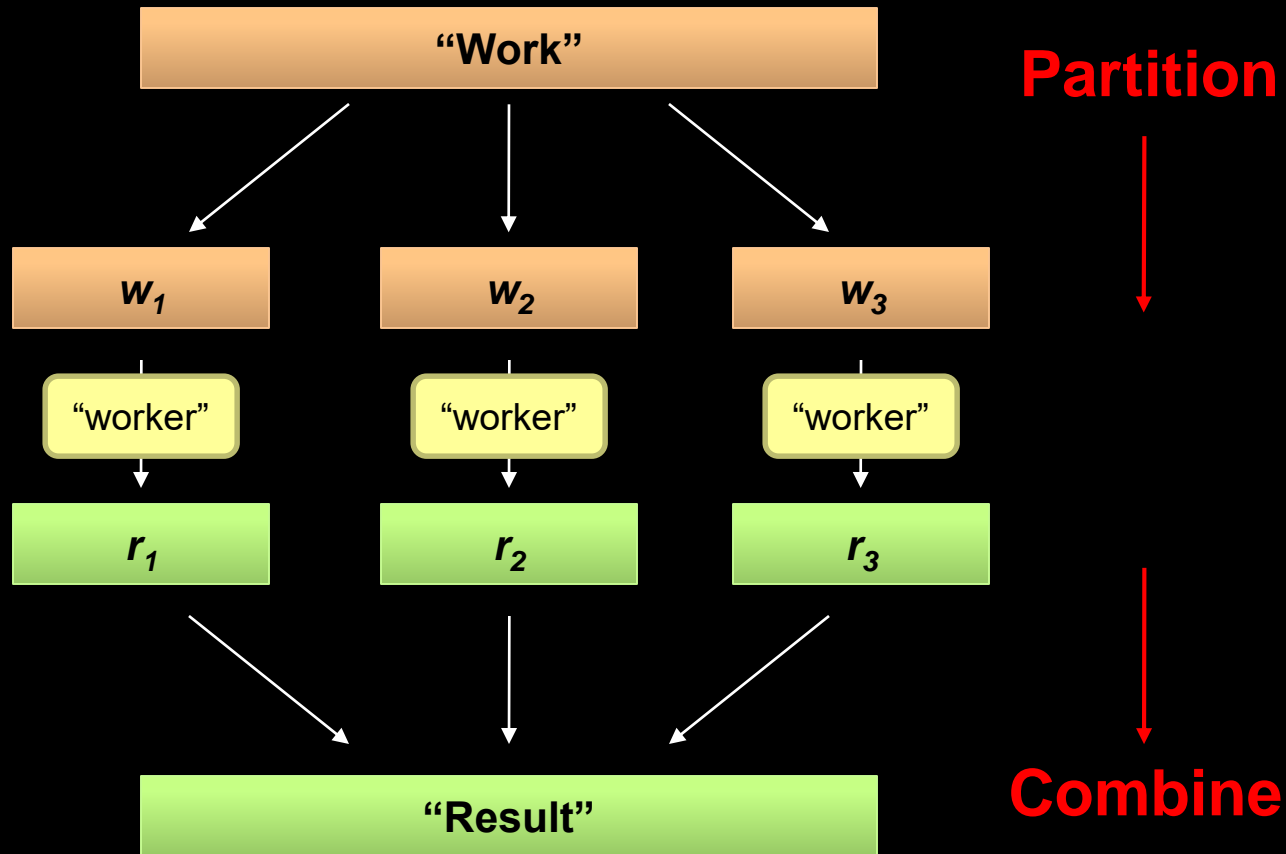
Ivory and SMRF

- Collaboration between:
 - University of Maryland
 - Yahoo! Research
- Reference implementation for a Web-scale IR toolkit
 - Designed around Hadoop from the ground up
 - Written specifically for the ClueWeb09 collection
 - Implements some of the algorithms described in this tutorial
 - Features SMRF query engine based on Markov Random Fields
- Open source
 - Initial release available now!

Cloud⁹

- Set of libraries originally developed for teaching MapReduce at the University of Maryland
 - Demos, exercises, etc.
- “Eat you own dog food”
 - Actively used for a variety of research projects


Divide and Conquer



It's a bit more complex...

Fundamental issues

scheduling, data distribution, synchronization,
inter-process communication, robustness, fault
tolerance, ...




Architectural issues

Flynn's taxonomy (SIMD, MIMD, etc.),
network topology, bisection bandwidth
UMA vs. NUMA, cache coherence

Common problems

livelock, deadlock, data starvation, priority inversion...
dining philosophers, sleeping barbers, cigarette smokers, ...

Different programming models



Different programming constructs

mutexes, conditional variables, barriers, ...
masters/slaves, producers/consumers, work queues, ...

**The reality: programmer shoulders the burden
of managing concurrency...**



Source: Ricardo Guimarães Herrmann

Typical Large-Data Problem

- Iterate over a large number of records

Map ○ Extract something of interest from each

- Shuffle and sort intermediate results

- Aggregate intermediate results

Reduce

- Generate final output

Key idea: provide a functional abstraction for these two operations

MapReduce Runtime

- Handles scheduling
 - Assigns workers to map and reduce tasks
- Handles “data distribution”
 - Moves processes to data
- Handles synchronization
 - Gathers, sorts, and shuffles intermediate data
- Handles faults
 - Detects worker failures and restarts
- Everything happens on top of a distributed FS (later)

“Hello World”: Word Count

Map(String docid, String text):

for each word w in text:

Emit(w, 1);

Reduce(String term, Iterator<Int> values):

int sum = 0;

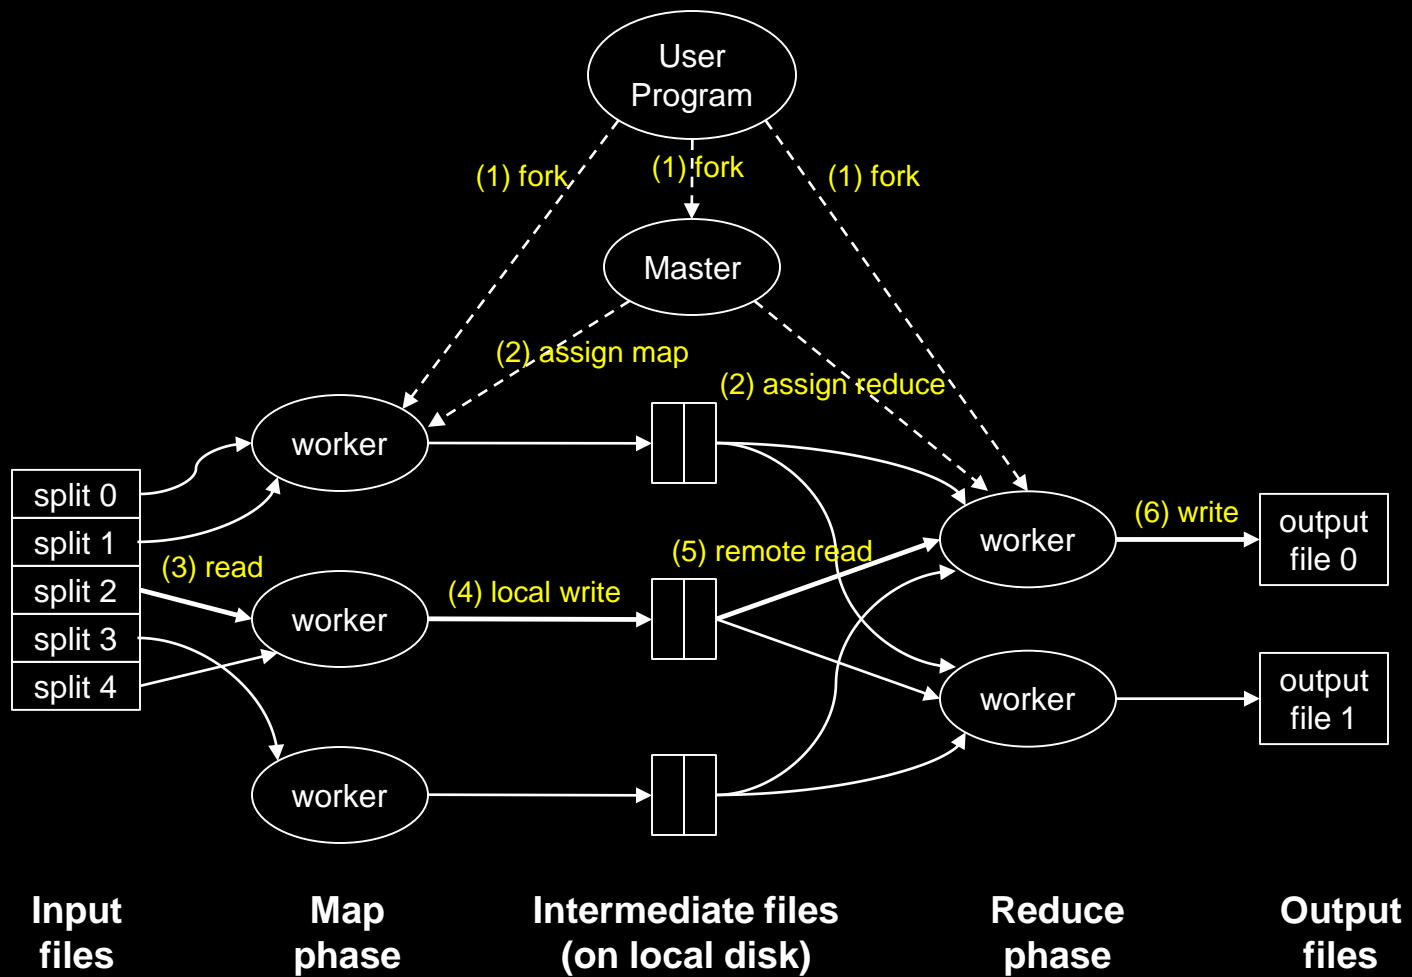
for each v in values:

sum += v;

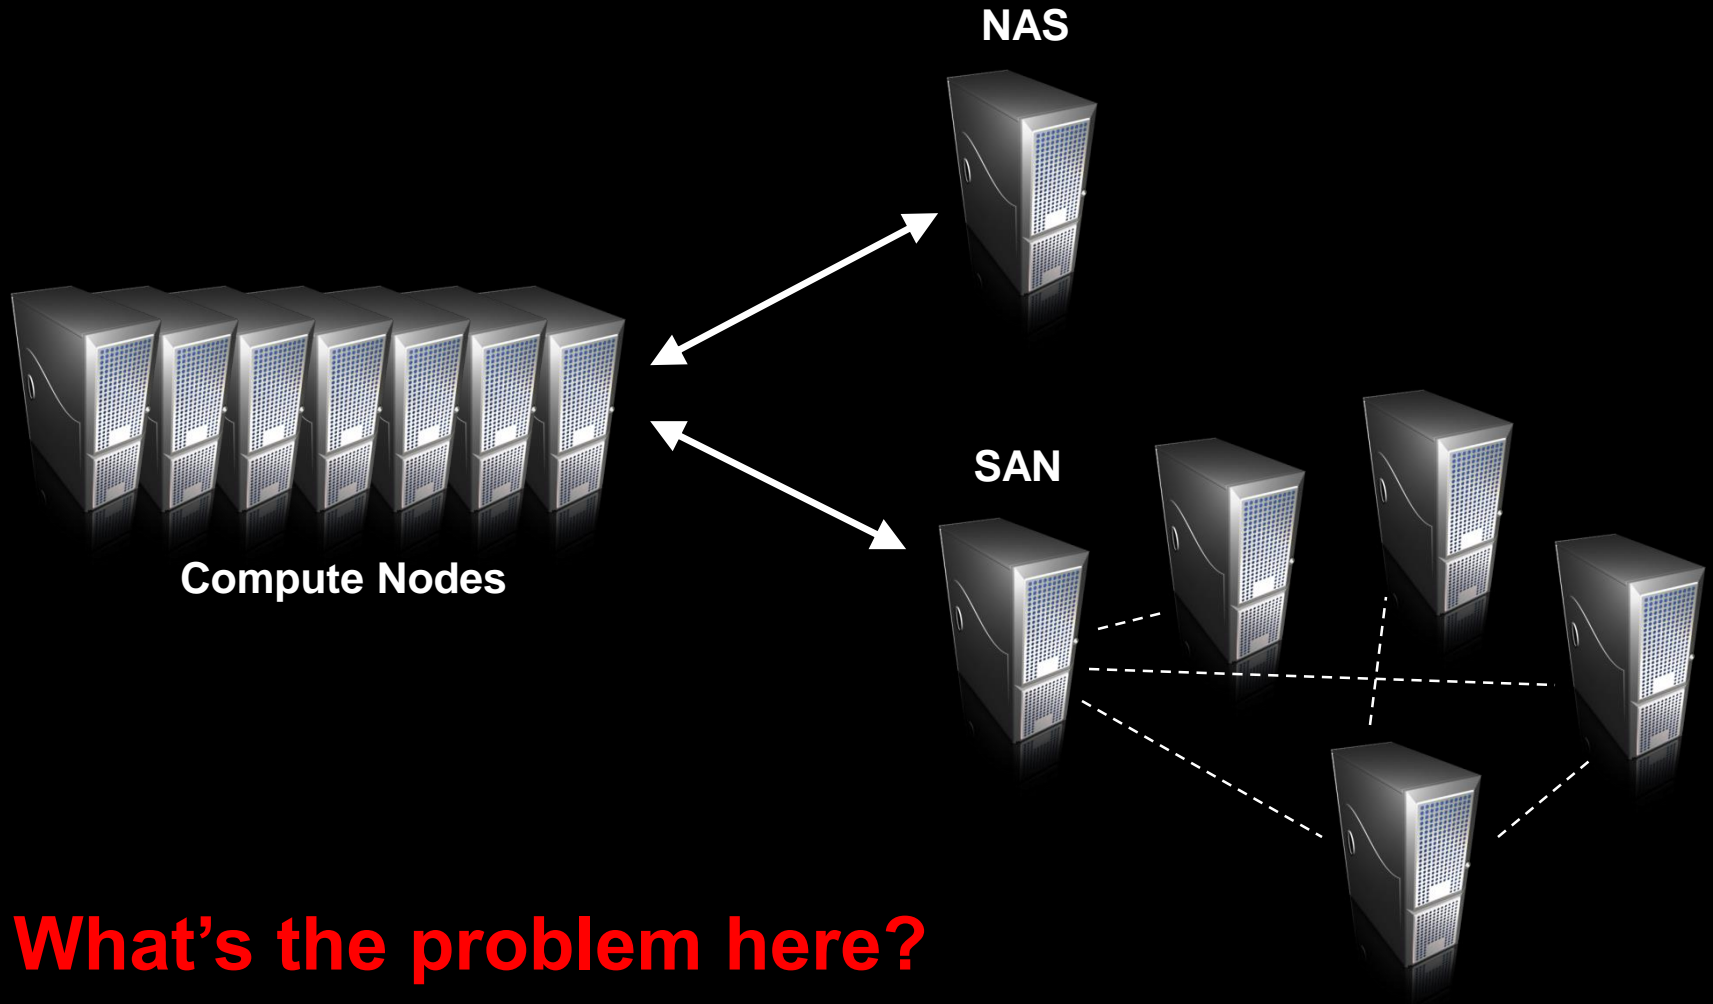
Emit(term, value);

MapReduce Implementations

- MapReduce is a programming model
- Google has a proprietary implementation in C++
 - Bindings in Java, Python
- Hadoop is an open-source implementation in Java
 - Project led by Yahoo, used in production
 - Rapidly expanding software ecosystem



How do we get data to the workers?



Distributed File System

- Don't move data to workers... move workers to the data!
 - Store data on the local disks of nodes in the cluster
 - Start up the workers on the node that has the data local
- Why?
 - Not enough RAM to hold all the data in memory
 - Disk access is slow, but disk throughput is reasonable
- A distributed file system is the answer
 - GFS (Google File System)
 - HDFS for Hadoop (= GFS clone)

GFS: Assumptions

- Commodity hardware over “exotic” hardware
 - Scale out, not up
- High component failure rates
 - Inexpensive commodity components fail all the time
- “Modest” number of HUGE files
- Files are write-once, mostly appended to
 - Perhaps concurrently
- Large streaming reads over random access
- High sustained throughput over low latency

GFS: Design Decisions

- Files stored as chunks
 - Fixed size (64MB)
- Reliability through replication
 - Each chunk replicated across 3+ chunkservers
- Single master to coordinate access, keep metadata
 - Simple centralized management
- No data caching
 - Little benefit due to large datasets, streaming reads
- Simplify the API
 - Push some of the issues onto the client

Master's Responsibilities

- Metadata storage
- Namespace management/locking
- Periodic communication with chunkservers
- Chunk creation, re-replication, rebalancing
- Garbage collection

Managing Dependencies

- Remember: Mappers run in isolation
 - You have no idea in what order the mappers run
 - You have no idea on what node the mappers run
 - You have no idea when each mapper finishes
- Tools for synchronization:
 - Ability to hold state in reducer across multiple key-value pairs
 - Sorting function for keys
 - Partitioner
 - Cleverly-constructed data structures

MapReduce: Large Counting Problems

- Term co-occurrence matrix for a text collection
= specific instance of a large counting problem
 - A large event space (number of terms)
 - A large number of observations (the collection itself)
 - Goal: keep track of interesting statistics about the events
- Basic approach
 - Mappers generate partial counts
 - Reducers aggregate partial counts

How do we aggregate partial counts efficiently?

First Try: “Pairs”

- Each mapper takes a sentence:
 - Generate all co-occurring term pairs
 - For all pairs, emit $(a, b) \rightarrow \text{count}$
- Reducers sums up counts associated with these pairs
- Use combiners!

“Pairs” Analysis

- Advantages
 - Easy to implement, easy to understand
- Disadvantages
 - Lots of pairs to sort and shuffle around (upper bound?)

Another Try: “Stripes”

- Idea: group together pairs into an associative array

$(a, b) \rightarrow 1$

$(a, c) \rightarrow 2$

$(a, d) \rightarrow 5$

$(a, e) \rightarrow 3$

$(a, f) \rightarrow 2$

$a \rightarrow \{ b: 1, c: 2, d: 5, e: 3, f: 2 \}$

- Each mapper takes a sentence:
 - Generate all co-occurring term pairs
 - For each term, emit $a \rightarrow \{ b: \text{count}_b, c: \text{count}_c, d: \text{count}_d \dots \}$
- Reducers perform element-wise sum of associative arrays

$$\begin{array}{r} a \rightarrow \{ b: 1, \quad d: 5, e: 3 \} \\ + \quad a \rightarrow \{ b: 1, c: 2, d: 2, \quad f: 2 \} \\ \hline a \rightarrow \{ b: 2, c: 2, d: 7, e: 3, f: 2 \} \end{array}$$

“Stripes” Analysis

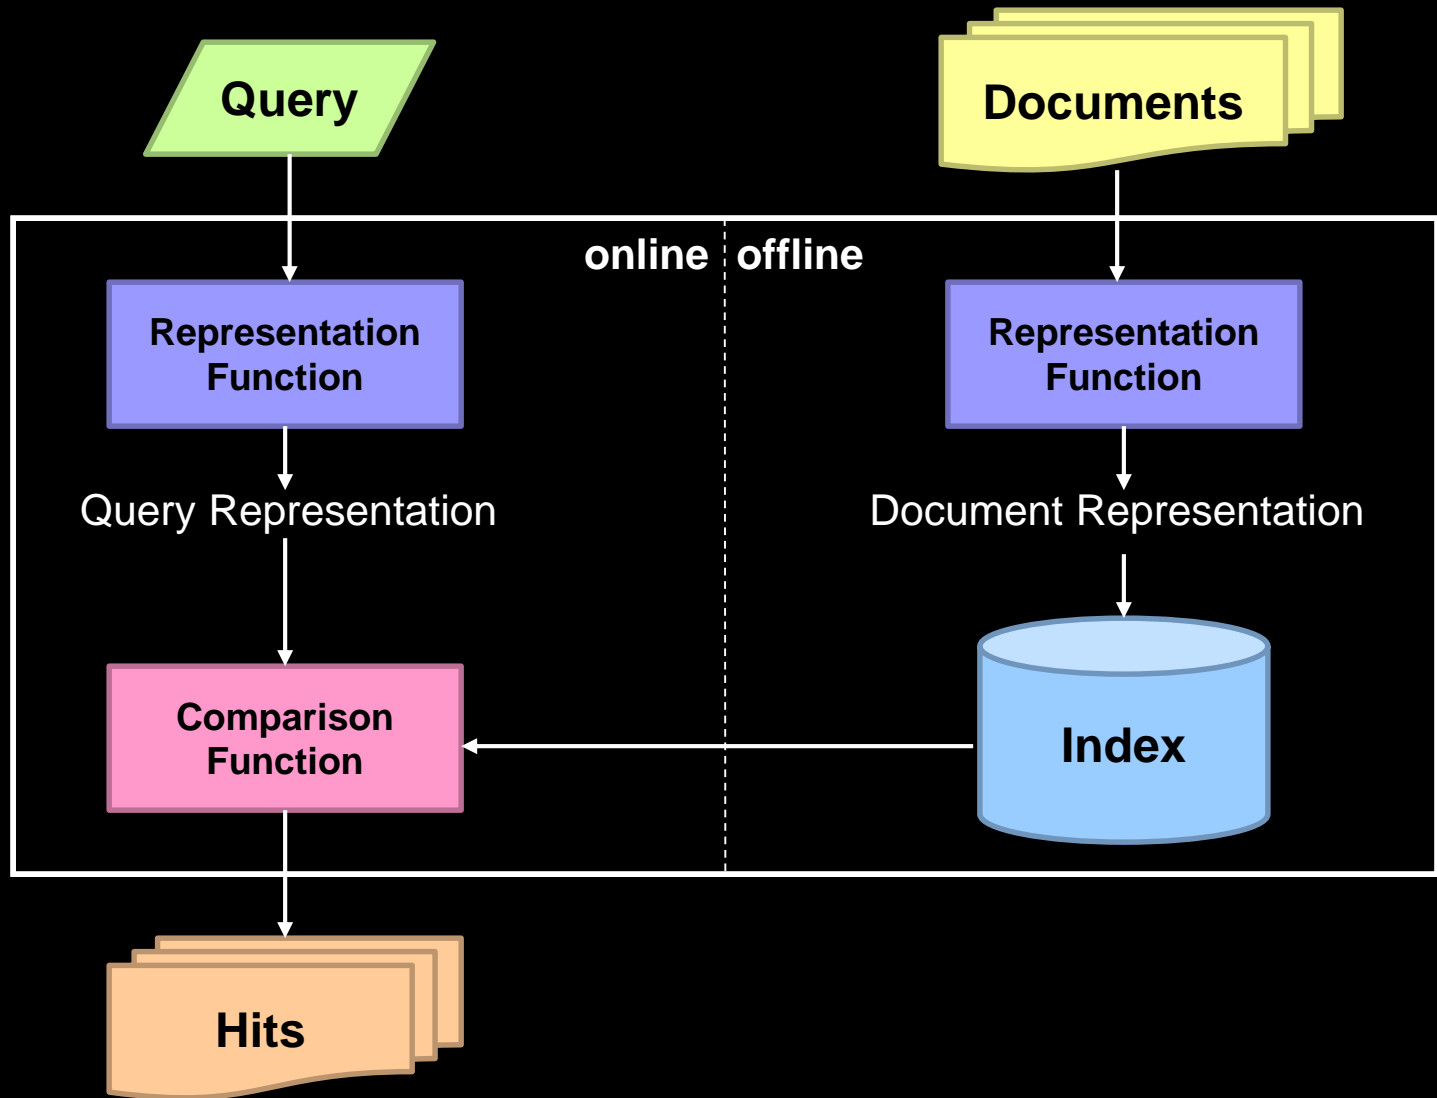
- Advantages

- Far less sorting and shuffling of key-value pairs
- Can make better use of combiners

- Disadvantages

- More difficult to implement
- Underlying object is more heavyweight
- Fundamental limitation in terms of size of event space

Abstract IR Architecture



MapReduce it?

- The indexing problem

- Scalability is critical
- Must be relatively fast, but need not be real time
- Fundamentally a batch operation
- Incremental updates may or may not be important
- For the web, crawling is a challenge in itself



Perfect for MapReduce!

- The retrieval problem

- Must have sub-second response time
- For the web, only need relatively few results

Uh... not so good...

Counting Words...



case folding, tokenization, stopwords removal, stemming

~~syntax~~, ~~semantics~~, ~~word knowledge~~, etc.

Inverted Index: Boolean Retrieval


Doc 1
one fish, two fish

Doc 2
red fish, blue fish

Doc 3
cat in the hat

Doc 4
green eggs and ham

	1	2	3	4
blue		1		
cat			1	
egg				1
fish	1	1		
green				1
ham				1
hat			1	
one	1			
red		1		
two	1			



blue	→	2
cat	→	3
egg	→	4
fish	→	1 2
green	→	4
ham	→	4
hat	→	3
one	→	1
red	→	2
two	→	1

Inverted Index: Ranked Retrieval


Doc 1
one fish, two fish

Doc 2
red fish, blue fish

Doc 3
cat in the hat

Doc 4
green eggs and ham

	<i>tf</i>				<i>df</i>
	1	2	3	4	
blue		1			1
cat			1		1
egg				1	1
fish	2	2			2
green				1	1
ham				1	1
hat			1		1
one	1				1
red		1			1
two	1				1



blue	→	1	→	2,1
cat	→	1	→	3,1
egg	→	1	→	4,1
fish	→	2	→	1,2 2,2
green	→	1	→	4,1
ham	→	1	→	4,1
hat	→	1	→	3,1
one	→	1	→	1,1
red	→	1	→	2,1
two	→	1	→	1,1

Inverted Index: Positional Information

Doc 1

one fish, two fish

Doc 2

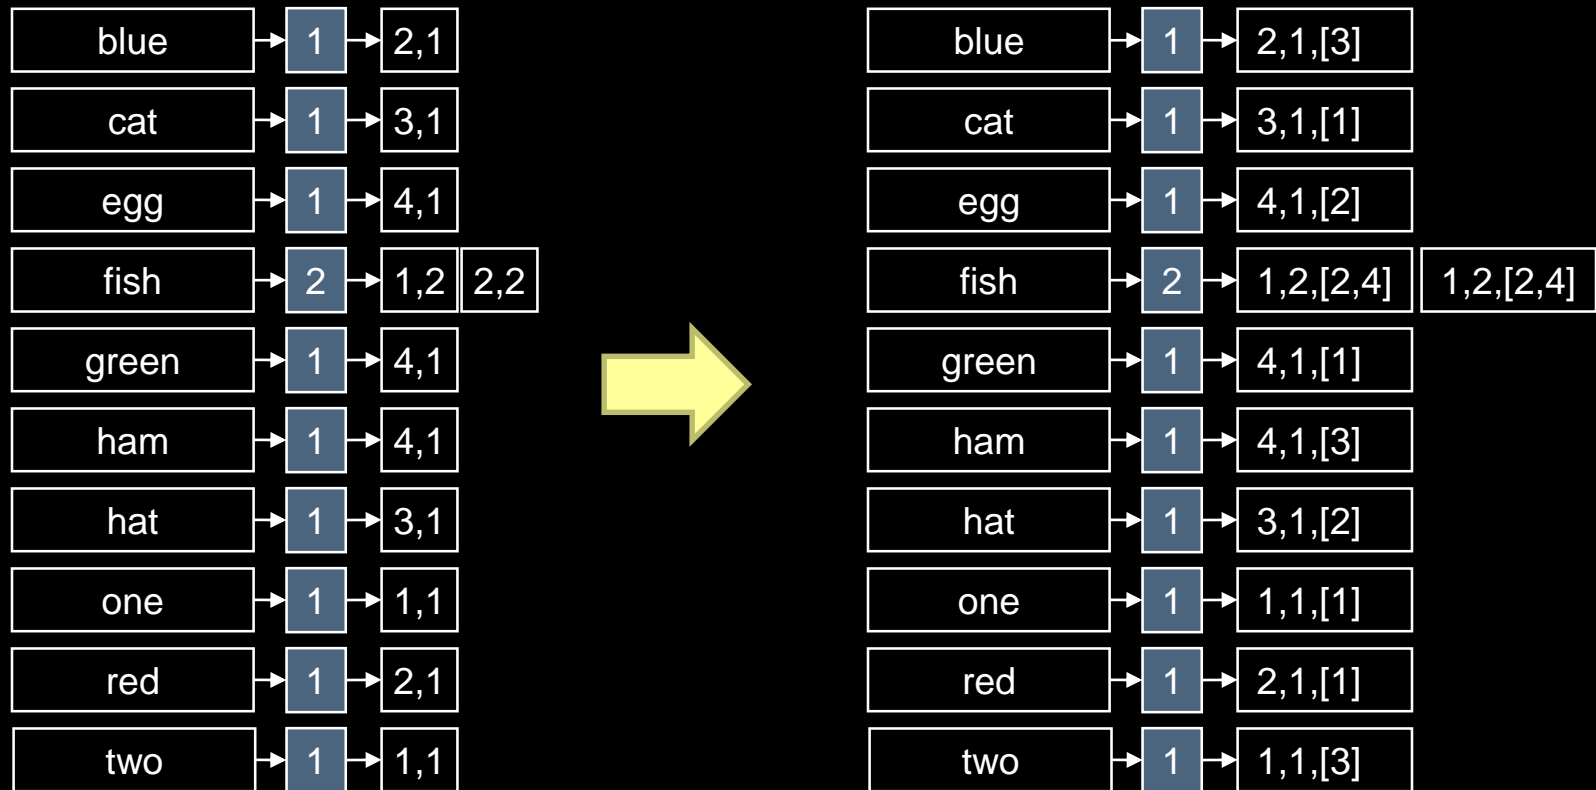
red fish, blue fish

Doc 3

cat in the hat

Doc 4

green eggs and ham



Indexing: Performance Analysis

- Fundamentally, a large sorting problem
 - Terms usually fit in memory
 - Postings usually don't
- How is it done on a single machine?
- How can it be done with MapReduce?
- First, let's characterize the problem size:
 - Size of vocabulary
 - Size of postings

MapReduce: Index Construction

- Map over all documents
 - Emit *term* as key, (*docno*, *tf*) as value
 - Emit other information as necessary (e.g., term position)
- Sort/shuffle: group postings by term
- Reduce
 - Gather and sort the postings (e.g., by *docno* or *tf*)
 - Write postings to disk
- MapReduce does all the heavy lifting!

Inverted Indexing with MapReduce

Map

Doc 1
one fish, two fish

one

1	1
---	---

two

1	1
---	---

fish

1	2
---	---

Doc 2
red fish, blue fish

red

2	1
---	---

blue

2	1
---	---

fish

2	2
---	---

Doc 3
cat in the hat

cat

3	1
---	---

hat

3	1
---	---

Shuffle and Sort: aggregate values by keys

Reduce

cat

3	1
---	---

fish

1	2
---	---

2	2
---	---

one

1	1
---	---

red

2	1
---	---

blue

2	1
---	---

hat

3	1
---	---

two

1	1
---	---

Inverted Indexing: Pseudo-Code

```
1: procedure MAP( $a, d$ )
2:   INITIALIZE.ASSOCIATIVEARRAY( $H$ )
3:   for all  $t \in d$  do
4:      $H\{t\} \leftarrow H\{t\} + 1$ 
5:   for all  $t \in H$  do
6:     EMIT( $t, \langle a, H\{t\} \rangle$ )

1: procedure REDUCE( $t, [\langle a_1, f_1 \rangle, \langle a_2, f_2 \rangle \dots]$ )
2:   INITIALIZE.LIST( $P$ )
3:   for all  $\langle a, f \rangle \in [\langle a_1, f_1 \rangle, \langle a_2, f_2 \rangle \dots]$  do
4:     APPEND( $P, \langle a, f \rangle$ )
5:   SORT( $P$ )
6:   EMIT( $t, P$ )
```

Positional Indexes

Map

Doc 1
one fish, two fish

one	1	1	[1]
two	1	1	[3]
fish	1	2	[2,4]

Doc 2
red fish, blue fish

red	2	1	[1]
blue	2	1	[3]
fish	2	2	[2,4]

Doc 3
cat in the hat

cat	3	1	[1]
hat	3	1	[2]

Shuffle and Sort: aggregate values by keys

Reduce

cat	3	1	[1]
fish	1	2	[2,4]
one	1	1	[1]
red	2	1	[1]

blue	2	1	[3]
hat	3	1	[2]
two	1	1	[3]

Inverted Indexing: Pseudo-Code

```
1: procedure MAP( $a, d$ )
2:   INITIALIZE.ASSOCIATIVEARRAY( $H$ )
3:   for all  $t \in d$  do
4:      $H\{t\} \leftarrow H\{t\} + 1$ 
5:   for all  $t \in H$  do
6:     EMIT( $t, \langle a, H\{t\} \rangle$ )

1: procedure REDUCE( $t, [\langle a_1, f_1 \rangle, \langle a_2, f_2 \rangle \dots]$ )
2:   INITIALIZE.LIST( $P$ )
3:   for all  $\langle a, f \rangle \in [\langle a_1, f_1 \rangle, \langle a_2, f_2 \rangle \dots]$  do
4:     APPEND( $P, \langle a, f \rangle$ )
5:   SORT( $P$ )
6:   EMIT( $t, P$ )
```


What's the problem?

Scalability Bottleneck

- Initial implementation: terms as keys, postings as values
 - Reducers must buffer all postings associated with key (to sort)
 - What if we run out of memory to buffer postings?
- Uh oh!

Another Try...

(key)	(values)		
fish	1	2	[2,4]
	34	1	[23]
	21	3	[1,8,22]
	35	2	[8,41]
	80	3	[2,9,76]
	9	1	[9]



	(keys)	(values)
fish	1	[2,4]
fish	9	[9]
fish	21	[1,8,22]
fish	34	[23]
fish	35	[8,41]
fish	80	[2,9,76]

How is this different?

- Let the framework do the sorting
- Term frequency implicitly stored
- Directly write postings to disk!

Wait, there's more!
(but first, an aside)

Postings Encoding

Conceptually:

fish

1	2
---	---

9	1
---	---

21	3
----	---

34	1
----	---

35	2
----	---

80	3
----	---

 ...

In Practice:

- Don't encode docnos, encode gaps (or d -gaps)
- But it's not obvious that this save space...

fish

1	2
---	---

8	1
---	---

12	3
----	---

13	1
----	---

1	2
---	---

45	3
----	---

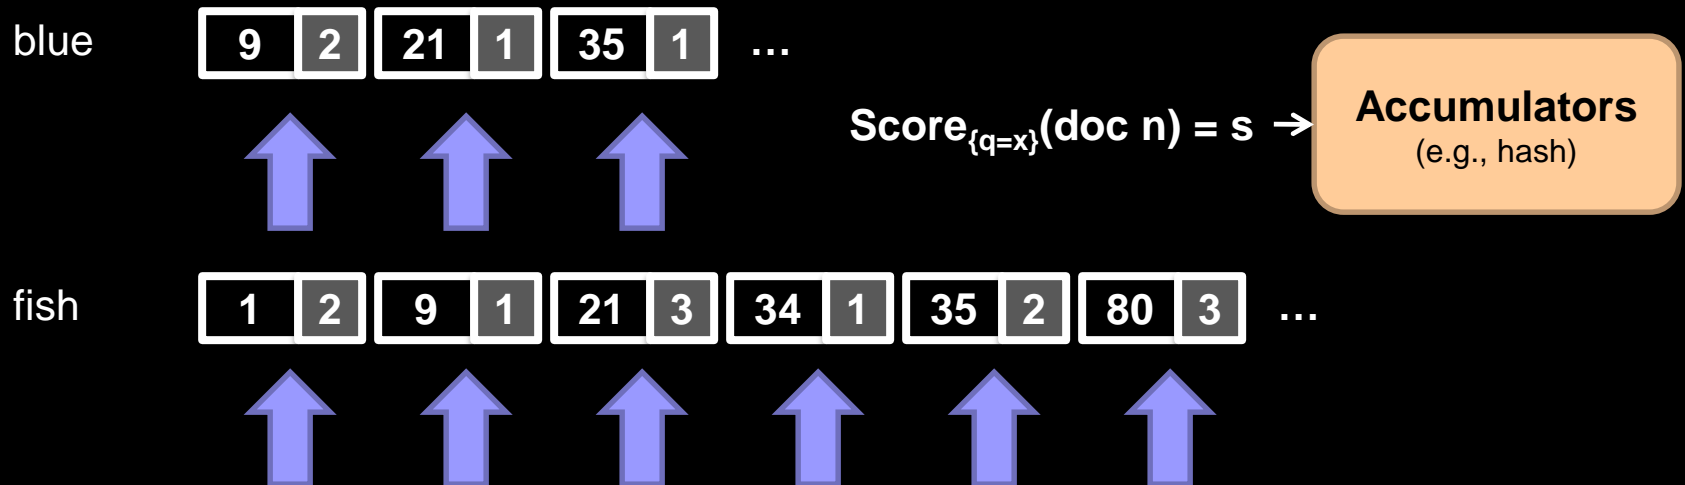
 ...

Retrieval in a Nutshell

- Look up postings lists corresponding to query terms
- Traverse postings for each query term
- Store partial query-document scores in *accumulators*
- Select top k results to return

Retrieval: Query-At-A-Time


- Evaluate documents one query at a time
 - Usually, starting from most rare term (often with tf-scored postings)



- Tradeoffs
 - Early termination heuristics (good)
 - Large memory footprint (bad), but filtering heuristics possible

Retrieval: Document-at-a-Time

- Evaluate documents one at a time (score all query terms)



Accumulators
(e.g. priority queue)

Document score in top k?

Yes: Insert document score, extract-min if queue too large

No: Do nothing

- Tradeoffs
 - Small memory footprint (good)
 - Must read through all postings (bad), but skipping possible
 - More disk seeks (bad), but blocking possible

Retrieval with MapReduce?

- MapReduce is fundamentally batch-oriented
 - Optimized for throughput, not latency
 - Startup of mappers and reducers is expensive
- MapReduce is not suitable for real-time queries!
 - Use separate infrastructure for retrieval...

Important Ideas

- Partitioning (for scalability)
- Replication (for redundancy)
- Caching (for speed)
- Routing (for load balancing)

The rest is just details!

When is MapReduce appropriate?

- Lots of input data
 - (e.g., compute statistics over large amounts of text)
 - Take advantage of distributed storage, data locality, aggregate disk throughput
- Lots of intermediate data
 - (e.g., postings)
 - Take advantage of sorting/shuffling, fault tolerance
- Lots of output data
 - (e.g., web crawls)
 - Avoid contention for shared resources
- Relatively little synchronization is necessary

When is MapReduce less appropriate?

- Data fits in memory
- Large amounts of shared data is necessary
- Fine-grained synchronization is needed
- Individual operations are processor-intensive